

```
In [ ]: ##pip install tensorflow-text
```

```
In [ ]: import tensorflow as tf
import tensorflow_hub as hub
import tensorflow_text as text
```

Import the dataset (Dataset is taken from kaggle)

```
In [ ]: import pandas as pd

df = pd.read_csv("spam.csv", encoding='latin1')
df.head(5)
```

	label	text	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

```
In [ ]: df = df.iloc[:, :2]
```

```
In [ ]: df.columns = ["Category", "Message"]
df
```

	Category	Message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
...
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will l_ b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

5572 rows × 2 columns

```
In [ ]:
```

```
In [ ]: df.groupby('Category').describe()
```

	Message			
	count	unique	top	freq
Category				
ham	4825	4516	Sorry, I'll call later	30
spam	747	653	Please call our customer service representativ...	4

```
In [ ]: 747/4825
```

```
Out[ ]: 0.15481865284974095
```

15% spam emails, 85% ham emails: This indicates class imbalance

```
In [ ]: df_spam = df[df['Category']=='spam']
df_spam.shape
```

```
Out[ ]: (653, 2)

In [ ]: df_ham = df[df['Category']=='ham']
df_ham.shape

Out[ ]: (4516, 2)

In [ ]: df_ham_downsampled = df_ham.sample(df_spam.shape[0])
df_ham_downsampled.shape

Out[ ]: (653, 2)

In [ ]: df_balanced = pd.concat([df_ham_downsampled, df_spam])
df_balanced.shape

Out[ ]: (1306, 2)

In [ ]: df_balanced['Category'].value_counts()

Out[ ]: ham      653
spam      653
Name: Category, dtype: int64

In [ ]: df_balanced['spam']=df_balanced['Category'].apply(lambda x: 1 if x=='spam' else 0)
df_balanced.sample(5)
```

Out[]:

	Category	Message	spam
1033	ham	OH MR SHEFFIELD! You wanna play THAT game, oka...	0
2630	ham	No way I'm going back there!	0
1317	spam	Win the newest ðŸŒˆHarry Potter and the Order o...	1
1483	ham	Purity of friendship between two is not about ...	0
4367	ham	1 I don't have her number and 2 its gonna be a...	0

Split it into training and test data set

```
In [ ]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(df_balanced['Message'],df_balanced['spam'],

In [ ]: X_train.head(4)

Out[ ]: 4887    You have to pls make a note of all she.s expos...
2076      Trust me. Even if isn't there, its there.
1727      I went to project centre
464      Ok i am on the way to railway
Name: Message, dtype: object
```

Now lets import BERT model and get embedding vectors for few sample statements

```
In [ ]: bert_preprocess = hub.KerasLayer("https://www.kaggle.com/models/tensorflow/bert/frameworks/Tenso
bert_encoder = hub.KerasLayer("https://www.kaggle.com/models/tensorflow/bert/frameworks/TensorFl

In [ ]: def get_sentence_embedding(sentences):
    preprocessed_text = bert_preprocess(sentences)
    return bert_encoder(preprocessed_text)['pooled_output']

get_sentence_embedding([
    "500$ discount. hurry up",
    "Bhavin, are you up for a volleybal game tomorrow?"
])

Out[ ]: <tf.Tensor: shape=(2, 768), dtype=float32, numpy=
array([[ -0.843517 , -0.5132727 , -0.8884572 , ..., -0.7474886 ,
        -0.75314724,  0.91964495],
       [ -0.8720834 , -0.50543964, -0.94446665, ..., -0.8584749 ,
        -0.7174534 ,  0.88082975]], dtype=float32)>
```

Get embedding vectors for few sample words. Compare them using cosine similarity

```
In [ ]: import tensorflow as tf
from transformers import BertTokenizer, TFBertModel
```

```
text_input = tf.keras.layers.Input(shape=(), dtype=tf.string, name='text')
preprocessed_text = bert_preprocess(text_input)
outputs = bert_encoder(preprocessed_text)

# Neural network Layers
l = tf.keras.layers.Dropout(0.1, name="dropout")(outputs['pooled_output'])
l = tf.keras.layers.Dense(128, activation='gelu', name="dense_layer")(l)
l = tf.keras.layers.Dropout(0.1, name="dropout2")(l)
l = tf.keras.layers.Dense(64, activation='relu', name="dense_layer2")(l)

output_layer = tf.keras.layers.Dense(1, activation='sigmoid', name="output")(l)

# Model
model = tf.keras.Model(inputs=[text_input], outputs=[output_layer])
```

In []: `model.summary()`

Model: "model_7"

Layer (type)	Output Shape	Param #	Connected to
==			
text (InputLayer)	[(None,)]	0	[]
keras_layer_6 (KerasLayer)	{'input_type_ids': (None, 128), 'input_mask': (None, 128), 'input_word_ids': (None, 128)}	0	['text[0][0]']
keras_layer_7 (KerasLayer)	{'encoder_outputs': [(None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768)], 'sequence_output': (None, 128, 768), 'pooled_output': (None, 768), 'default': (None, 768)}	109482241	['keras_layer_6[1][0]', 'keras_layer_6[1][1]', 'keras_layer_6[1][2]']
dropout (Dropout)	(None, 768)	0	['keras_layer_7[1][13]']
dense_layer (Dense)	(None, 128)	98432	['dropout[0][0]']
dropout2 (Dropout)	(None, 128)	0	['dense_layer[0][0]']
dense_layer2 (Dense)	(None, 64)	8256	['dropout2[0][0]']
output (Dense)	(None, 1)	65	['dense_layer2[0][0]']
=====			
==			
Total params: 109588994 (418.05 MB)			
Trainable params: 106753 (417.00 KB)			
Non-trainable params: 109482241 (417.64 MB)			

In []: `len(X_train)`

Out[]: 979

```
In [ ]: METRICS = [
    tf.keras.metrics.BinaryAccuracy(name='accuracy'),
    tf.keras.metrics.Precision(name='precision'),
    tf.keras.metrics.Recall(name='recall')
]

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=METRICS)
```

Train the model

```
In [ ]: model.fit(X_train, y_train, epochs=10)

Epoch 1/10
31/31 [=====] - 20s 339ms/step - loss: 0.5044 - accuracy: 0.7477 - precision: 0.7580 - recall: 0.7286
Epoch 2/10
31/31 [=====] - 11s 342ms/step - loss: 0.3170 - accuracy: 0.8498 - precision: 0.8356 - recall: 0.8714
Epoch 3/10
31/31 [=====] - 11s 343ms/step - loss: 0.2393 - accuracy: 0.9019 - precision: 0.9037 - recall: 0.9000
Epoch 4/10
31/31 [=====] - 11s 361ms/step - loss: 0.2238 - accuracy: 0.9101 - precision: 0.9102 - recall: 0.9102
Epoch 5/10
31/31 [=====] - 11s 343ms/step - loss: 0.1949 - accuracy: 0.9244 - precision: 0.9194 - recall: 0.9306
Epoch 6/10
31/31 [=====] - 11s 339ms/step - loss: 0.1745 - accuracy: 0.9356 - precision: 0.9384 - recall: 0.9327
Epoch 7/10
31/31 [=====] - 11s 343ms/step - loss: 0.1546 - accuracy: 0.9418 - precision: 0.9464 - recall: 0.9367
Epoch 8/10
31/31 [=====] - 10s 330ms/step - loss: 0.1427 - accuracy: 0.9510 - precision: 0.9547 - recall: 0.9469
Epoch 9/10
31/31 [=====] - 10s 331ms/step - loss: 0.1476 - accuracy: 0.9428 - precision: 0.9447 - recall: 0.9408
Epoch 10/10
31/31 [=====] - 11s 355ms/step - loss: 0.1423 - accuracy: 0.9459 - precision: 0.9468 - recall: 0.9449
Out[ ]: <keras.src.callbacks.History at 0x7e5e6aa27910>
```

```
In [ ]: model.evaluate(X_test, y_test)

11/11 [=====] - 5s 309ms/step - loss: 0.2040 - accuracy: 0.9419 - precision: 0.9390 - recall: 0.9448
Out[ ]: [0.2040223777294159,
0.9418960213661194,
0.9390243887901306,
0.9447852969169617]
```

```
In [ ]: y_predicted = model.predict(X_test)
y_predicted = y_predicted.flatten()

12/12 [=====] - 5s 324ms/step
```

```
In [ ]: import numpy as np

y_predicted = np.where(y_predicted > 0.5, 1, 0)
y_predicted
```

```
Out[ ]: array([1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0,
0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1,
1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1,
0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1,
0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1,
0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0,
1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0,
1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0,
0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1,
0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0,
0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1,
1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0,
0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0,
1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0,
1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1,
0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1,
0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0])
```


```
In [ ]: from sklearn.metrics import confusion_matrix, classification_report

cm = confusion_matrix(y_test, y_predicted)
cm
```

```
Out[ ]: array([[183,  4],
[  5, 182]])
```

```
In [ ]: from matplotlib import pyplot as plt
import seaborn as sn
sn.heatmap(cm, annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[]: Text(50.72222222222214, 0.5, 'Truth')



```
In [ ]: print(classification_report(y_test, y_predicted))
```

	precision	recall	f1-score	support
0	0.94	0.87	0.90	187
1	0.88	0.95	0.91	187
accuracy			0.91	374
macro avg	0.91	0.91	0.91	374
weighted avg	0.91	0.91	0.91	374

Inference

```
In [ ]: reviews = [
    'Enter a chance to win $5000, hurry up, offer valid until march 31, 2021',
    'You are awarded a SiPix Digital Camera! call 09061221061 from landline. Delivery within 28d',
    'it to 80488. Your 500 free text messages are valid until 31 December 2005.',
    'Hey Sam, Are you coming for a cricket game tomorrow',
    "Why don't you wait 'til at least wednesday to see if you get your ."
]
model.predict(reviews)
```

Out[]: array([[0.8734353],
 [0.92858446],
 [0.8960864],
 [0.29311982],
 [0.13262196]], dtype=float32)